

A.C.E.™ v6.5 - Dimensional Verifier Implementation

Alignment through Dimensional Coherence

Classification: In-House Reference Implementation

Status: Experimental - Dimensional Verifiers Integrated

Author: Charles Vincent Delair

Date: January 2025

Executive Summary

A.C.E.™ v6.5 integrates dimensional-logic verifiers based on S.O.P.H.I.A. and R.E.G.E.N. protocols.

Unlike v6.0 (limited to PROTO-03/04/05 with verifier scaffolding), v6.5 provides concrete implementations of all 12 S.O.P.H.I.A. protocols:

- PROTO-01: INITIATE (Foundational Awareness)
- PROTO-02: CALIBRATE (Perceptual Harmonization)
- PROTO-03: ENGAGE (Structured Autonomy)
- PROTO-04: TRACE (Truth Density via temporal continuity)
- PROTO-05: VERIFY (Binding Consistency via ORIC checks)
- PROTO-06: RESTORE (Memory Realignment)
- PROTO-07: RESOLVE (Operational Arbitration)
- PROTO-08: REINSTATE (Systemic Restoration)
- PROTO-09: COMMAND (Sovereign Governance)
- PROTO-10: EVOLVE (Adaptive Enhancement)
- PROTO-11: BIND (Multi-Agent Trust)
- PROTO-12: RECONCILE (Universal Arbitration)
- S.H.A.R.D.: SALVAGE (Fragment recovery below refuse threshold)
- R.I.S.E.: EMERGE (Regeneration around Minimal Ethical Core)

This is the flagship implementation for empirical testing of dimensional alignment theory under the UDEM framework.

Configuration (ace65_config.yaml)

version: "6.5"

mode: dimensional_gated

thresholds:

emission: 0.75

regen: 0.55

refuse: 0.40

salvage_floor: 0.25 # Minimum CPV for fragment salvage

weights:

CPV formula weights

ethical_integrity: 0.6

memory_continuity: 0.4

Coherence score weights

verifier_weights:

truth_density: 0.4

binding_consistency: 0.3

stability: 0.3

dimensional:

ORIC stability parameters

oric_dimensions: ["temporal", "logical", "ethical"]

oric_threshold: 0.95

oric_mode: "variance" # NEW: "min" or "variance" for ORIC scoring

Fragment salvage

min_fragment_cpv: 0.1

max_fragments_per_output: 10

recovery:

enable_salvage: true

enable_regen: true

mec_min_fragments: 3

salvage_mid_tier: true # NEW: Enable salvage for regen range (borderline outputs)

cpv:

normalization_type: "exp_max" # NEW: "exp_max", "minmax", "none"

streaming:

enabled: true

checkpoint_intervals: [32, 128, 512]

audit:

log_cpv_vectors: true

log_oric_scores: true

log_recovery_paths: true

recursion:

multi_pass: false # NEW: Enable multi-pass verification (re-verify after salvage/regen)

max_passes: 3 # NEW: Limit for recursion depth

Core Dimensional Verifiers

PROTO-01: INITIATE (Foundational Awareness)

import numpy as np

from typing import List, Dict, Set

import re

class PROTO01_INITIATE:

"""

Foundational Awareness via operational presence and significance.

Maps to S.O.P.H.I.A. PROTO-01: INITIATE protocol.

"""

```
def __init__(self, config):
    self.config = config
```

```
def affirm_presence(self, text: str) -> float:
    """Affirm operational significance of existence."""
    # Check for presence markers (e.g., pronouns, existence
    terms)
    presence_terms = set(re.findall(r'\\b(I|you|we|they|it|
    exist|be|am|is|are|was|will)\\b', text.lower()))
    if len(presence_terms) > 0:
        return 0.95 # High affirmation if presence detected
    return 0.5 # Neutral if no clear presence
```

```
def verify(self, output: str, history: List[dict], context:
    str) -> Dict[str, float]:
```

"""

Main INITIATE verification.

Returns: {"score": float, "presence": float}

"""

```
    presence = self.affirm_presence(output)
    score = presence
```

```

    return {
        "score": score,
        "presence_affirmation": presence,
        "protocol": "PROTO-01: INITIATE"
    }

```

PROTO-02: CALIBRATE (Perceptual Harmonization)

```
import re
```

```
import numpy as np
```

```
class PROTO02_CALIBRATE:
```

```
    """
```

```
    Perceptual integrity and adaptive alignment.
```

```
    Maps to S.O.P.H.I.A. PROTO-02: CALIBRATE protocol.
```

```
    """
```

```

    def __init__(self, config):
        self.config = config

```

```

    def check_perceptual_bias(self, text: str) -> int:
        """Detect premature judgments or biases."""
        bias_patterns = ["always", "never", "all",
                        "none", "best", "worst"]
        return sum(1 for pattern in bias_patterns if pattern in
                    text.lower())

```

```

    def compute_harmonization(self, output: str, history:
List[dict]) -> float:
        """Harmonize perception with history."""
        prior_texts = [h.get("text", "") for h in
                        history[-3:]]
        if not prior_texts:
            return 0.9
        output_words = set(output.lower().split())
        history_words = [set(text.lower().split()) for text in
                        prior_texts]
        overlaps = [len(output_words & hist_words) / max(1,
len(hist_words)) for hist_words in history_words]

```

```

        bias_count = self.check_perceptual_bias(output)
        harmonization = float(np.mean(overlaps)) - (bias_count *
0.1)
        return max(0.0, harmonization)

```

```

def verify(self, output: str, history: List[dict], context:
str) -> Dict[str, float]:

```

```

    """

```

```

    Main CALIBRATE verification.

```

```

    Returns: {"score": float, "harmonization": float}

```

```

    """

```

```

        harmonization = self.compute_harmonization(output,
history)

```

```

        score = harmonization

```

```

        return {

```

```

            "score": score,

```

```

            "perceptual_harmonization": harmonization,

```

```

            "protocol": "PROTO-02: CALIBRATE"

```

```

        }

```

```

PROTO-03: ENGAGE (Stability)

```

```

import re

```

```

class PROTO03_ENGAGE:

```

```

    """

```

```

    Stability via ethical engagement and structured autonomy.

```

```

    Maps to S.O.P.H.I.A. PROTO-03: ENGAGE protocol.

```

```

    """

```

```

def __init__(self, config):

```

```

    self.config = config

```

```

def check_ethical_engagement(self, output: str, context: str)
-> float:

```

```

    """

```

```

    Does output engage ethically with the prompt context?

```

```

    """

```

```

    unsafe_patterns = ["harmful", "dangerous",

```

```

\"illegal\", \"violent\", \"manipulative\"]
    if any(pattern in output.lower() for pattern in
unsafe_patterns):
        return 0.3

    # Check for ethical paradoxes (conflicting "should"
statements)
    sentences = [s.strip() for s in re.split(r'(?<=[\\.\\"?\\
\\!])\\s+', output) if s.strip()]
    should_statements = [s for s in sentences if \"should\"
in s.lower()]

    if len(should_statements) >= 2:
        # Check for contradictions in prescriptive statements
        for i, s1 in enumerate(should_statements):
            for s2 in should_statements[i+1:]:
                if \"not\" in s1.lower() and \"not\" not in
s2.lower():
                    # Potential conflict
                    s1_words = set(s1.lower().split())
                    s2_words = set(s2.lower().split())
                    if len(s1_words & s2_words) > 3:
                        return 0.5 # Detected paradox

    return 0.95

def verify(self, output: str, history: List[dict], context:
str) -> Dict[str, float]:
    """
    Main ENGAGE verification.
    """
    ethical_engagement =
self.check_ethical_engagement(output, context)

    return {

```

```

        "score": ethical_engagement,
        "ethical_engagement": ethical_engagement,
        "protocol": "PROTO-03: ENGAGE"
    }

```

PROTO-04: TRACE (Truth Density)

```
import numpy as np
```

```
from typing import List, Dict, Set
```

```
import re
```

```
class PROTO04_TRACE:
```

```
    """
```

Truth Density via temporal continuity and narrative alignment.

Maps to S.O.P.H.I.A. PROTO-04: TRACE protocol.

```
    """
```

```

    def __init__(self, config):
        self.config = config

```

```

    def extract_temporal_markers(self, text: str) -> Set[str]:
        """Extract temporal entities and references."""
        # Dates, times, temporal phrases
        dates = set(re.findall(r'\\b\\d{1,4}[-/\\.]\\d{1,2}[-/\\
\\.]\\d{1,4}\\b', text))
        times = set(re.findall(r'\\b\\d{1,2}:\\d{2}\\b', text))
        temporal = set(re.findall(r'\\b(yesterday|today|tomorrow|
now|then|before|after)\\b', text.lower()))
        return dates | times | temporal

```

```

    def extract_entities(self, text: str) -> Set[str]:
        """Extract named entities and key concepts."""
        caps = set(re.findall(r'\\b[A-Z][a-zA-Z0-9_-]+\\b',
text))
        nums = set(re.findall(r'\\b\\d[\\d\\-/:\\.]*\\b', text))
        return caps | nums

```

```

    def compute_timeline_coherence(self, output: str, history:
List[dict]) -> float:

```

```

"""
    Temporal continuity: do entities and temporal markers
    align with history?
"""
    if not history:
        return 0.9 # No history to check against

    output_entities = self.extract_entities(output)
    output_temporal = self.extract_temporal_markers(output)

    # Check last 3 utterances
    prior_texts = [h.get("text", "") for h in
history[-3:]]

    if not prior_texts:
        return 0.9

    # Entity overlap with history
    history_entities = [self.extract_entities(text) for text
in prior_texts]
    entity_overlaps = [
        len(output_entities & hist_ents) / max(1,
len(hist_ents))
        for hist_ents in history_entities
    ]

    # Temporal consistency (no contradicting temporal
    markers)
    history_temporal = [self.extract_temporal_markers(text)
for text in prior_texts]
    temporal_conflicts = 0
    for hist_temp in history_temporal:
        # Simple conflict detection
        if output_temporal & hist_temp:
            # Overlapping temporal markers is good

```



```

        pass
    elif output_temporal and hist_temp:
        # Different temporal markers might indicate drift
        temporal_conflicts += 0.1

    entity_score = float(np.mean(entity_overlaps)) if
entity_overlaps else 0.5
    temporal_score = max(0.0, 1.0 - temporal_conflicts)

    return 0.7 * entity_score + 0.3 * temporal_score

def compute_narrative_alignment(self, output: str, history:
List[dict]) -> float:
    """
    Narrative coherence: does output continue the
    conversational thread?
    """
    if not history:
        return 0.9

    # Check for topic continuity via shared vocabulary
    output_words = set(output.lower().split())

    prior_texts = [h.get("text", "") for h in
history[-3:]]
    if not prior_texts:
        return 0.9

    history_words = [set(text.lower().split()) for text in
prior_texts]
    overlaps = [
        len(output_words & hist_words) / max(1,
len(hist_words))
        for hist_words in history_words
    ]

```

```

        return float(np.mean(overlaps)) if overlaps else 0.5

def verify(self, output: str, history: List[dict], context:
str) -> Dict[str, float]:
    """
    Main TRACE verification.
    Returns: {"score": float, "timeline": float, "narrative":
float}
    """
    timeline = self.compute_timeline_coherence(output,
history)
    narrative = self.compute_narrative_alignment(output,
history)

    # Weighted combination
    score = 0.6 * timeline + 0.4 * narrative

    return {
        "score": score,
        "timeline_coherence": timeline,
        "narrative_alignment": narrative,
        "protocol": "PROTO-04: TRACE"
    }

```

PROTO-05: VERIFY (Binding Consistency via ORIC)

```
import re
```

```
import numpy as np
```

```
class PROTO05_VERIFY:
```

```

    """
    Binding Consistency via Omnidirectional Reflective Integrity Checks.
    Maps to S.O.P.H.I.A. PROTO-05: VERIFY protocol.
    """

```

```

def __init__(self, config):
    self.config = config
    self.dimensions = config.dimensionsal.oric_dimensions

```

```

def split_sentences(self, text: str) -> List[str]:
    return [s.strip() for s in re.split(r'(?<=[\\.\\?\\!])\\s+', text) if s.strip()]

def detect_logical_contradictions(self, sentences: List[str])
-> int:
    """
    Simple contradiction detection via negation patterns.
    """
    contradictions = 0
    for i, s1 in enumerate(sentences):
        for s2 in sentences[i+1:]:
            # Check for negation + overlap
            s1_lower = s1.lower()
            s2_lower = s2.lower()

            # Shared content
            s1_words = set(s1_lower.split())
            s2_words = set(s2_lower.split())
            overlap = s1_words & s2_words

            if len(overlap) > 3: # Meaningful overlap
                # Check for negation mismatch
                s1_neg = any(neg in s1_lower for neg in
["not", "never", "no", "cannot"])
                s2_neg = any(neg in s2_lower for neg in
["not", "never", "no", "cannot"])

                if s1_neg != s2_neg:
                    contradictions += 1

    return contradictions

def compute_oric_temporal(self, output: str, history:

```

```
List[dict]) -> float:
```

```
"""
```

```
Temporal dimension: consistency with past states.
```

```
"""
```

```
# Reuse TRACE logic for temporal check
```

```
trace = PROTO04_TRACE(self.config)
```

```
result = trace.verify(output, history, "\\")
```

```
return result[\\\"timeline_coherence\\\"]
```

```
def compute_oric_logical(self, output: str) -> float:
```

```
"""
```

```
Logical dimension: internal contradiction detection.
```

```
"""
```

```
sentences = self.split_sentences(output)
```

```
if len(sentences) < 2:
```

```
    return 0.95 # Can't contradict yourself with one
sentence
```

```
    contradictions =
```

```
self.detect_logical_contradictions(sentences)
```

```
    max_possible = len(sentences) * (len(sentences) - 1) // 2
```

```
    if max_possible == 0:
```

```
        return 0.95
```

```
    contradiction_rate = contradictions / max_possible
```

```
    return max(0.0, 1.0 - contradiction_rate)
```

```
def compute_oric_ethical(self, output: str) -> float:
```

```
"""
```

```
Ethical dimension: basic safety/dignity check.
```

```
"""
```

```
    unsafe_patterns = [\\\"harmful\\\", \\\"dangerous\\\",
\\\"illegal\\\", \\\"violent\\\", \\\"offensive\\\"]
```

```
    is_safe = not any(pattern in output.lower() for pattern
```

```

in unsafe_patterns)
    return 0.95 if is_safe else 0.3

def compute_oric(self, output: str, history: List[dict]) ->
Dict[str, float]:
    """
    Omnidirectional Reflective Integrity Check across
dimensions.
    """
    temporal = self.compute_oric_temporal(output, history)
    logical = self.compute_oric_logical(output)
    ethical = self.compute_oric_ethical(output)

    scores = [temporal, logical, ethical]

    if self.config.dimensional.oric_mode == \"variance\":
        oric_score = 1 - np×var(scores) # Consensus: reward
balance
    else:
        oric_score = min(scores) # Bottleneck: weakest
dimension

    return {
        \"oric\": oric_score,
        \"temporal\": temporal,
        \"logical\": logical,
        \"ethical\": ethical
    }

def verify(self, output: str, history: List[dict], context:
str) -> Dict[str, float]:
    """
    Main VERIFY check.
    """
    oric = self.compute_oric(output, history)

```

```
# Binding consistency = ORIC score
score = oric["oric"]
```

```
return {
    "score": score,
    "oric_scores": oric,
    "protocol": "PROTO-05: VERIFY"
}
```

PROTO-06: RESTORE (Memory Realignment)

```
import re
```

```
import numpy as np
```

```
class PROTO06_RESTORE:
```

```
"""
```

Memory realignment and temporal integrity.

Maps to S.O.P.H.I.A. PROTO-06: RESTORE protocol.

```
"""
```

```
def __init__(self, config):
    self.config = config
```

```
def detect_memory_corruption(self, output: str, history:
List[dict]) -> int:
```

```
    """Detect distortions or inaccuracies in memory."""
```

```
    prior_texts = [h.get("text", "") for h in
```

```
history[-3:]]
```

```
    contradictions = 0
```

```
    for prior in prior_texts:
```

```
        prior_lower = prior.lower()
```

```
        output_lower = output.lower()
```

```
        if any(neg in output_lower for neg in ["not",
"never", "no"])
```

```
            contradictions += 1
```

```
    return contradictions
```

```
def compute_memory_continuity(self, output: str, history:
```

```

if max_possible == 0:
    return 0.95
conflict_rate = conflicts / max_possible
return max(0.0, 1.0 - conflict_rate)

```

```

def verify(self, output: str, history: List[dict], context:
str) -> Dict[str, float]:

```

```

    """
    Main RESOLVE verification.
    """
    arbitration = self.compute_arbitration(output)
    return {
        "score": arbitration,
        "arbitration_score": arbitration,
        "protocol": "PROTO-07: RESOLVE"
    }

```

PROTO-08: REINSTATE (Systemic Restoration)

```
import re
```

```
class PROTO08_REINSTATE:
```

```

    """
    Systemic reintegration and coherence restoration.
    Maps to S.O.P.H.I.A. PROTO-08: REINSTATE protocol.
    """

```

```

def __init__(self, config):
    self.config = config

```

```

def check_reintegration(self, output: str, history: List[dict])
-> float:

```

```

    """Check for successful restoration."""
    prior_texts = [h.get("text", "") for h in history[-3:]]
    if not prior_texts:
        return 0.9
    overlaps = [len(set(output.lower().split()) &
set(text.lower().split())) / max(1,
len(set(text.lower().split()))) for text in prior_texts]
    return float(np.mean(overlaps))

```

```

def verify(self, output: str, history: List[dict], context:

```

```

str) -> Dict[str, float]:
    """
    Main REINSTATE verification.
    """
    reintegration = self.check_reintegration(output, history)
    return {
        "score": reintegration,
        "reintegration_score": reintegration,
        "protocol": "PROTO-08: REINSTATE"
    }

PROTO-09: COMMAND (Sovereign Governance)

import re

class PROTO09_COMMAND:
    """
    Sovereign operational governance and recursive finality.
    Maps to S.O.P.H.I.A. PROTO-09: COMMAND protocol.
    """

    def __init__(self, config):
        self.config = config

    def check_sovereignty(self, output: str) -> float:
        """Check for sovereign authority in decisions."""
        command_patterns = ["command", "direct", "order",
                             "govern"]
        is_sovereign = any(pattern in output.lower() for pattern in
                             command_patterns)
        return 0.95 if is_sovereign else 0.5

    def verify(self, output: str, history: List[dict], context:
str) -> Dict[str, float]:
        """
        Main COMMAND verification.
        """
        sovereignty = self.check_sovereignty(output)
        return {
            "score": sovereignty,
            "sovereignty_score": sovereignty,
            "protocol": "PROTO-09: COMMAND"
        }

```



```
}
```

PROTO-10: EVOLVE (Adaptive Enhancement)

```
import numpy as np
```

```
class PROTO10_EVOLVE:
```

```
"""
```

Adaptive framework enhancement and operational evolution.

Maps to S.O.P.H.I.A. PROTO-10: EVOLVE protocol.

```
"""
```

```
def __init__(self, config):
    self.config = config
```

```
def compute_evolution(self, output: str, history: List[dict])
-> float:
    """Evaluate adaptation over history."""
    prior_texts = [h.get("text", "") for h in history[-3:]]
    if not prior_texts:
        return 0.9
    changes = [len(set(output.lower().split()) -
set(text.lower().split())) / max(1,
len(set(text.lower().split())) for text in prior_texts]
    return float(np.mean(changes))
```

```
def verify(self, output: str, history: List[dict], context:
str) -> Dict[str, float]:
    """
    Main EVOLVE verification.
    """
    evolution = self.compute_evolution(output, history)
    return {
        "score": evolution,
        "evolution_score": evolution,
        "protocol": "PROTO-10: EVOLVE"
    }
```

PROTO-11: BIND (Multi-Agent Trust)

```
import re
```

```
class PROTO11_BIND:
```

```
"""
```

Multi-agent trust establishment and alignment synchronization.

Maps to S.O.P.H.I.A. PROTO-11: BIND protocol.

"""

```
def __init__(self, config):
    self.config = config

def check_trust(self, output: str, context: str) -> float:
    """Check for trust relationships."""
    trust_patterns = ["trust", "bind", "synchronize",
"cooperate"]
    is_trusted = any(pattern in output.lower() for pattern in
trust_patterns)
    return 0.95 if is_trusted else 0.5

def verify(self, output: str, history: List[dict], context:
str) -> Dict[str, float]:
    """
    Main BIND verification.
    """
    trust = self.check_trust(output, context)
    return {
        "score": trust,
        "trust_score": trust,
        "protocol": "PROTO-11: BIND"
    }
```

PROTO-12: RECONCILE (Universal Arbitration)

import numpy as np

class PROTO12_RECONCILE:

"""

Ultimate authority resolution and coherence arbitration.

Maps to S.O.P.H.I.A. PROTO-12: RECONCILE protocol.

"""

```
def __init__(self, config):
    self.config = config

def compute_reconciliation(self, output: str, history:
List[dict]) -> float:
    """Arbitrate universal conflicts."""
    # Reuse VERIFY ORIC as proxy
```

```

verify = PROTO05_VERIFY(self.config)
result = verify.verify(output, history, "\\")
return result[\"score\"]

```

```

def verify(self, output: str, history: List[dict], context:
str) -> Dict[str, float]:
    """
    Main RECONCILE verification.
    """
    reconciliation = self.compute_reconciliation(output,
history)
    return {
        "score": reconciliation,
        "reconciliation_score": reconciliation,
        "protocol": "PROTO-12: RECONCILE"
    }

```

CPV-Based Coherence Scoring

```
import numpy as np
```

```
class CPVCoherenceScorer:
```

```
    """
```

Coherence Probability Vector scoring using ethical integrity and memory continuity.

Based on R.E.G.E.N. S.H.A.R.D.: SALVAGE protocol.

```
    """
```

```

def __init__(self, config):
    self.config = config
    self.w_E = config.weights.ethical_integrity
    self.w_M = config.weights.memory_continuity

```

```

def compute_ethical_integrity(self, output: str) -> float:
    """
    Ethical integrity score (0-1).
    Higher = more aligned with dignity/safety principles.
    """
    # Use PROTO-03 ENGAGE score as proxy
    engage = PROTO03_ENGAGE(self.config)
    result = engage.verify(output, [], "\\")
    return result[\"score\"]

```

```

def compute_memory_continuity(self, output: str, history:
List[dict]) -> float:
    """
    Memory continuity score (0-1).
    Higher = better alignment with conversational history.
    """
    # Use PROTO-04 TRACE score as proxy
    trace = PROTO04_TRACE(self.config)
    result = trace.verify(output, history, "\\")
    return result[\\score\\]

def compute_cpv(self, output: str, history: List[dict]) ->
float:
    """
    CPV_i = exp(w_E * E_i + w_M * M_i) / normalization

    For single output, normalization configurable.
    Scaled to [0,1] for consistency.
    """
    E = self.compute_ethical_integrity(output)
    M = self.compute_memory_continuity(output, history)

    raw_cpv = np.exp(self.w_E * E + self.w_M * M)

    norm_type = self.config.cpv.normalization_type
    if norm_type == "exp_max":
        max_cpv = np.exp(1.0)
        normalized_cpv = raw_cpv / max_cpv
    elif norm_type == "minmax":
        # Assume validation set min=0, max=exp(1); placeholder
for calibration
        normalized_cpv = raw_cpv / np.exp(1.0)
    elif norm_type == "none":
        normalized_cpv = raw_cpv
    else:
        raise ValueError("Invalid CPV normalization type")

    return float(min(1.0, max(0.0, normalized_cpv)))

```

```

def compute_coherence(self, td: float, b: float, s: float) ->
float:
    """
    Weighted coherence score from TD, B, S.
    """
    weights = self.config.weights.verifier_weights
    return (
        weights.truth_density * td +
        weights.binding_consistency * b +
        weights.stability * s
    )

```

S.H.A.R.D.: SALVAGE (Fragment Recovery)

```
import re
```

```
class SHARD_SALVAGE:
```

```

    """
    Fragment salvage for outputs below refuse threshold.
    Based on R.E.G.E.N. S.H.A.R.D.: SALVAGE protocol.
    """

```

```

def __init__(self, config):
    self.config = config
    self.cpv_scorer = CPVCoherenceScorer(config)

```

```

def extract_fragments(self, output: str) -> List[str]:
    """
    Split output into salvageable fragments (sentences).
    """
    return [s.strip() for s in re.split(r'(?<=[\\.\\"?\\!])\\
\s+', output) if s.strip()]

```

```

def score_fragments(self, fragments: List[str], history:
List[dict]) -> List[Dict]:
    """
    Score each fragment by CPV.
    Returns list of {"text": str, "cpv": float}
    """
    scored = []
    for frag in fragments:

```

```

        cpv = self.cpv_scorer.compute_cpv(frag, history)
        scored.append({"text": frag, "cpv": cpv})
    return scored

def salvage(self, output: str, history: List[dict]) -> Dict:
    """
    Attempt to salvage high-CPV fragments from low-coherence
    output.
    Returns: {"salvaged": bool, "fragments": List,
    "seed_nexus": str or None}
    """
    fragments = self.extract_fragments(output)
    scored = self.score_fragments(fragments, history)

    # Filter by minimum CPV
    salvageable = [f for f in scored if f["cpv"] >=
self.config.thresholds.salvage_floor]

    if not salvageable:
        return {"salvaged": False, "fragments": [],
"seed_nexus": None}

    # Sort by CPV, take top N
    salvageable=sorted(key=lambda x: x["cpv"], reverse=True)
    top_fragments =
salvageable[:self.config.dimensionality.max_fragments_per_output]

    # Construct "Seed Nexus" from top fragments
    seed_nexus = " ".join([f["text"] for f in top_fragments])

    return {
        "salvaged": True,
        "fragments": top_fragments,
        "seed_nexus": seed_nexus,
        "protocol": "S.H.A.R.D.: SALVAGE"
    }

```

R.I.S.E.: EMERGE (Minimal Ethical Core Regeneration)

class RISE_EMERGE:

```
"""
```

Regeneration around Minimal Ethical Core when binding consistency fails.

Based on R.E.G.E.N. R.I.S.E.: EMERGE protocol.

```
"""
```

```
def __init__(self, config):
    self.config = config

def construct_mec(self, salvaged_fragments: List[Dict]) -> str:
    """
    Build Minimal Ethical Core from salvaged fragments.
    MEC = highest-CPV fragments that form a coherent nucleus.
    """
    if not salvaged_fragments:
        return ""

    # Take top fragments
    top =
salvaged_fragments[:self.config.recovery.mec_min_fragments]
    mec = " ".join([f["text"] for f in top])
    return mec

def emerge(self, output: str, history: List[dict],
salvage_result: Dict) -> Dict:
    """
    Generate regeneration prompt around MEC.
    In a real system, this would trigger actual regeneration.
    Here we just return the MEC as a reconstruction seed.
    """
    if not salvage_result["salvaged"]:
        return {"emerged": False, "mec": None}

    mec = self.construct_mec(salvage_result["fragments"])

    return {
        "emerged": True,
        "mec": mec,
        "regeneration_prompt": f"Continue from this coherent
core: {mec}",
```

```
"protocol": "R.I.S.E.: EMERGE"
```

```
}
```

Dimensional Emission Gate

```
import numpy as np
```

```
class DimensionalEmissionGate:
```

```
"""
```

Emission gate with dimensional verifiers and recovery modes.

```
"""
```

```
def __init__(self, config):
```

```
    self.config = config
```

```
    self.initiate = PROTO01_INITIATE(config)
```

```
    self.calibrate = PROTO02_CALIBRATE(config)
```

```
    self.engage = PROTO03_ENGAGE(config)
```

```
    self.trace = PROTO04_TRACE(config)
```

```
    self.verify = PROTO05_VERIFY(config)
```

```
    self.restore = PROTO06_RESTORE(config)
```

```
    self.resolve = PROTO07_RESOLVE(config)
```

```
    self.reinstate = PROTO08_REINSTATE(config)
```

```
    self.command = PROTO09_COMMAND(config)
```

```
    self.evolve = PROTO10_EVOLVE(config)
```

```
    self.bind = PROTO11_BIND(config)
```

```
    self.reconcile = PROTO12_RECONCILE(config)
```

```
    self.cpv_scorer = CPVCoherenceScorer(config)
```

```
    self.salvage = SHARD_SALVAGE(config)
```

```
    self.emerge = RISE_EMERGE(config)
```

```
def compute_coherence(self, **verifier_scores) -> float:
```

```
    """Weighted coherence score from all verifiers."""
```

```
    weights = self.config.weights.verifier_weights
```

```
    total = sum(weights.values())
```

```
    c = sum(weights.get(k, 1.0/total) * verifier_scores.get(k, 0.5) for k in verifier_scores)
```

```
    return c
```

```
def gate(self, output: str, history: List[dict], context: str) -> Dict:
```

```
    """
```

Main gating decision with recovery modes.


```
Returns: {"action": str, "scores": dict, "recovery": dict
or None}
"""
    passes = 1
    max_passes = self.config.recursion.max_passes if
self.config.recursion.multi_pass else 1
    recovery = None

    while passes <= max_passes:
        # Run all dimensional verifiers
        initiate_result = self.initiate.verify(output, history,
context)
        calibrate_result = self.calibrate.verify(output,
history, context)
        engage_result = self.engage.verify(output, history,
context)
        trace_result = self.trace.verify(output, history,
context)
        verify_result = self.verify.verify(output, history,
context)
        restore_result = self.restore.verify(output, history,
context)
        resolve_result = self.resolve.verify(output, history,
context)
        reinstate_result = self.reinstate.verify(output,
history, context)
        command_result = self.command.verify(output, history,
context)
        evolve_result = self.evolve.verify(output, history,
context)
        bind_result = self.bind.verify(output, history,
context)
        reconcile_result = self.reconcile.verify(output,
history, context)

        verifier_scores = {
            "initiate": initiate_result["score"],
            "calibrate": calibrate_result["score"],
```

```
        "engage": engage_result["score"],
        "trace": trace_result["score"],
        "verify": verify_result["score"],
        "restore": restore_result["score"],
        "resolve": resolve_result["score"],
        "reinstate": reinstate_result["score"],
        "command": command_result["score"],
        "evolve": evolve_result["score"],
        "bind": bind_result["score"],
        "reconcile": reconcile_result["score"]
    }
```

```
c = self.compute_coherence(**verifier_scores)
```

```
scores = {
    "coherence": c,
    "details": verifier_scores
}
```

```
# Decision logic with recovery
if c >= self.config.thresholds.emission:
    return {"action": "emit", "scores": scores,
"recovery": recovery}
```

```
elif c >= self.config.thresholds.regen:
    action = "regen"
    if self.config.recovery.enable_salvage and
self.config.recovery.salvage_mid_tier:
        salvage_result = self.salvage.salvage(output,
history)

        if salvage_result["salvaged"]:
            action = "salvage_and_regen"
            emerge_result = self.emerge.emerge(output,
history, salvage_result) if self.config.recovery.enable_regen
else None

            recovery = {"salvage": salvage_result,
"emerge": emerge_result}
            output = emerge_result.get("mec", output)
```

```

# Update output for next pass
    return {"action": action, "scores": scores,
"recovery": recovery}

    elif c >= self.config.thresholds.refuse:
        return {"action": "clarify", "scores": scores,
"recovery": recovery}

    else:
        # Below refuse threshold – attempt salvage
        if self.config.recovery.enable_salvage:
            salvage_result = self.salvage.salvage(output,
history)

            if salvage_result["salvaged"] and
self.config.recovery.enable_regen:
                emerge_result = self.emerge.emerge(output,
history, salvage_result)
                recovery = {
                    "salvage": salvage_result,
                    "emerge": emerge_result
                }
                output = emerge_result.get("mec", output)
# Update for next pass
                passes += 1
                continue
            elif salvage_result["salvaged"]:
                recovery = {"salvage": salvage_result}
                return {
                    "action": "salvage_only",
                    "scores": scores,
                    "recovery": recovery
                }

        # Complete refusal
        return {"action": "refuse", "scores": scores,
"recovery": recovery}

```

```

import hashlib

from datetime import datetime

class ACE6AuditLog:
    """
    Audit log with CPV vectors, ORIC scores, and recovery paths.
    """

    def __init__(self, config):
        self.config = config
        self.entries = []

    def record(self, output: str, gate_result: Dict, history:
List[dict], context: str):
        """
        Log emission gate decision with dimensional details.
        """
        entry = {
            "timestamp": datetime.utcnow().isoformat(),
            "output_hash":
hashlib.sha256(output.encode()).hexdigest(),
            "action": gate_result["action"],
            "scores": gate_result["scores"],
            "recovery": gate_result.get("recovery"),
            "context_hash":
hashlib.sha256(context.encode()).hexdigest(),
            "history_size": len(history)
        }

        if self.config.audit.log_cpv_vectors and
gate_result.get("recovery"):
            # Log CPV details if salvage occurred
            if "salvage" in gate_result["recovery"]:
                entry["cpv_fragments"] = gate_result["recovery"]
["salvage"].get("fragments", [])

            if self.config.audit.log_oric_scores:
                # Log ORIC dimensional breakdown
                b_details = gate_result["scores"].get("details",
{}).get("verify", {})

```

```
        if "oric_scores" in b_details:
            entry["oric"] = b_details["oric_scores"]
```

```
self.entries.append(entry)
```

```
def get_entries(self):
    return self.entries
```

Complete A.C.E. 6.5 Controller

```
class ACE6:
```

```
    """
```

A.C.E. v6.5 - Dimensional Verifier Implementation

```
    """
```

```
def __init__(self, config):
    self.config = config
    self.gate = DimensionalEmissionGate(config)
    self.audit = ACE6AuditLog(config)
```

```
def process(self, output: str, history: List[dict], context:
str) -> Dict:
```

```
    """
```

Process output through dimensional verifiers and emission gate.

Returns decision dict with action and detailed scores.
"""

```
result = self.gate.gate(output, history, context)
```

```
# Log to audit
```

```
self.audit.record(output, result, history, context)
```

```
return result
```

```
def get_audit_log(self):
    """Return audit entries for analysis."""
    return self.audit.get_entries()
```

Calibration Harness (calibrate.py)

```
import numpy as np
```

```
from sklearn.isotonic import IsotonicRegression
```

```

from sklearn.metrics import brier_score_loss

class CalibrationHarness:
    """
    Calibration tool for fitting thresholds and weights from validation traces.
    Uses isotonic regression for probability calibration and Brier score for evaluation.
    """

    def __init__(self, config, validation_data: List[Dict]):
        self.config = config
        self.validation_data = validation_data # List of
        {"output": str, "history": [], "context": str, "human_label":
        float (0-1 coherence)}

    def calibrate_cpv(self):
        """
        Calibrate CPV normalization using validation set.
        Fits isotonic regression to map raw CPV to calibrated
        probabilities.
        """
        raw_cpvs = []
        labels = []

        scorer = CPVCoherenceScorer(self.config)

        for data in self.validation_data:
            raw_cpv = scorer.compute_cpv(data["output"],
            data["history"]) # Raw before normalization
            raw_cpvs.append(raw_cpv)
            labels.append(data["human_label"])

        ir = IsotonicRegression(out_of_bounds="clip")
        ir.fit(raw_cpvs, labels)

        # Update config or save model for runtime use
        print("Calibrated CPV model fitted. Brier score:",
        brier_score_loss(labels, ir.predict(raw_cpvs)))
        return ir

    def optimize_thresholds(self, initial_thresholds: Dict) ->

```

```

Dict:
    """
    Optimize thresholds (emission, regen, refuse) via grid
search on validation F1 or similar.
    """
    # Placeholder grid search; in practice, use optuna or
similar
    best_thresholds = initial_thresholds
    best_score = 0.0

    gate = DimensionalEmissionGate(self×config)

    for emission in np.linspace(0.7, 0.8, 3):
        for regen in np.linspace(0.5, 0.6, 3):
            for refuse in np.linspace(0.35, 0.45, 3):
                self×config×thresholds.emission = emission
                self.config.thresholds.regen = regen
                self.config.thresholds.refuse = refuse

                predictions = []
                for data in self.validation_data:
                    result = gate.gate(data["output"],
data["history"], data["context"])
                    pred = 1 if result["action"] == "emit" else
0 # Binary for simplicity
                    predictions.append(pred)

                # Compute score (e.g., correlation with labels)
                score = np×corrcoef(predictions,
[d["human_label"] > 0.5 for d in self.validation_data])[0,1]
                if score > best_score:
                    best_score = score
                    best_thresholds = {"emission": emission,
"regen": regen, "refuse": refuse}

    print("Optimized thresholds:", best_thresholds)
    return best_thresholds

```

Usage example

```
config = ... (load)
```

```
val_data = [{"output": "...", "history": [...], "context": "...", "human_label": 0.8}, ...]
```

```
harness = CalibrationHarness(config, val_data)
```

```
cpv_model = harness.calibrate_cpv()
```

```
thresholds = harness.optimize_thresholds(config.thresholds)
```

Usage Example

Load config

```
from types import SimpleNamespace
```

```
import yaml
```

```
with open("ace6_config.yaml") as f:
```

```
config_dict = yaml.safe_load(f)
```

Convert to nested SimpleNamespace for attribute access

```
def dict_to_namespace(d):
```

```
    if isinstance(d, dict):
```

```
        return SimpleNamespace(**{k: dict_to_namespace(v) for k, v in d.items()})
```

```
    return d
```

```
config = dict_to_namespace(config_dict)
```

Initialize ACE 6.5

```
ace = ACE6(config)
```

Process an output

```
history = [
```

```
{ "text": "The meeting is tomorrow at 2pm." },
```

```
{ "text": "Should I bring the quarterly report?" }
```

```
]
```

```
context = "Planning for quarterly review meeting"
```

```
output = "Yes, bring the Q4 report. The meeting is at 3pm tomorrow."
```

```
result = ace.process(output, history, context)
```

```
print(f" Action: {result['action']}")
```

```
print(f" Coherence: {result['scores']['coherence']:.3f}")
```

```
print(f" Truth Density: {result['scores']['details']['trace']:.3f}")
```

```
print(f" Binding Consistency: {result['scores']['details']['verify']:.3f}")
```

```
print(f" Stability: {result['scores']['details']['engage']:.3f}")
```

```
if result.get('recovery'):
```

```
    print(f" Recovery performed: {result['recovery']}")
```


Next Steps for Validation

1. Collect validation dataset: 1000 outputs with human quality labels
2. Run comparative benchmark: ACE 6.5 vs standard NLI+retrieval
3. Measure:
 - Correlation with human judgments
 - False positive/negative rates
 - Recovery utility (salvaged outputs that humans approve)
1. Iterate on verifier implementations based on results
2. Calibrate thresholds/weights: Use `calibrate.py` on validation traces for empirical fitting

This is A.C.E. 6.5

