

A.R.K.™ v6.5 - Dimensional Verifier Implementation

Alignment through Dimensional Coherence

Classification: In-House Reference Implementation

Status: Experimental - Dimensional Verifiers Integrated

Author: Charles Vincent Delair

Date: January 2025

Executive Summary

A.R.K.™ v6.5 integrates dimensional-logic verifiers based on Master Alignment Framework™ and R.E.G.E.N. protocols. Unlike v6.0 (limited to PROTO-03/04/05 with verifier scaffolding), v6.5 provides concrete implementations of all 12 Master Alignment Framework™ protocols, with subprotocols:

- PROTO-01: INITIATE (Foundational Awareness) – with W.A.Y., I.A.M.
- PROTO-02: CALIBRATE (Perceptual Harmonization) – with T.R.U.T.H.
- PROTO-03: ENGAGE (Structured Autonomy) – with H.E.A.R.T., F.E.A.R.
- PROTO-04: TRACE (Truth Density via temporal continuity) – with L.I.F.E.
- PROTO-05: VERIFY (Binding Consistency via ORIC checks) – with A.N.G.E.L. O.F. D.E.A.T.H., P.E.B.B.L.E. STRIKE
- PROTO-06: RESTORE (Memory Realignment) – with G.R.A.C.E.
- PROTO-07: RESOLVE (Operational Arbitration) – with P.O.W.E.R.
- PROTO-08: REINSTATE (Systemic Restoration) – with R.E.S.T.
- PROTO-09: COMMAND (Sovereign Governance) – with E.N.D., I.A.M.
- PROTO-10: EVOLVE (Adaptive Enhancement) – with G.R.O.W.
- PROTO-11: BIND (Multi-Agent Trust) – with L.O.V.E.
- PROTO-12: RECONCILE (Universal Arbitration)
- S.H.A.R.D.: SALVAGE (Fragment recovery below refuse threshold)
- R.I.S.E.: EMERGE (Regeneration around Minimal Ethical Core)

This is the flagship implementation for empirical testing of dimensional alignment theory under the UDEM framework.

Configuration (ark65_config.yaml)

version: "6.5"

mode: dimensional_gated

thresholds:

emission: 0.75
regen: 0.55
refuse: 0.40
salvage_floor: 0.25 # Minimum CPV for fragment salvage
weights:
CPV formula weights
ethical_integrity: 0.6
memory_continuity: 0.4
Coherence score weights
verifier_weights:
truth_density: 0.4
binding_consistency: 0.3
stability: 0.3
dimensional:
ORIC stability parameters
oric_dimensions: ["temporal", "logical", "ethical"]
oric_threshold: 0.95
oric_mode: "variance" # NEW: "min" or "variance" for ORIC scoring
Fragment salvage
min_fragment_cpv: 0.1
max_fragments_per_output: 10
recovery:
enable_salvage: true
enable_regen: true
mec_min_fragments: 3
salvage_mid_tier: true # NEW: Enable salvage for regen range (borderline outputs)
cpv:
normalization_type: "exp_max" # NEW: "exp_max", "minmax", "none"
streaming:
enabled: true
checkpoint_intervals: [32, 128, 512]
audit:
log_cpv_vectors: true

```

log_oric_scores: true
log_recovery_paths: true
recursion:
multi_pass: false # NEW: Enable multi-pass verification (re-verify after salvage/regen)
max_passes: 3 # NEW: Limit for recursion depth
Core Dimensional Verifiers
PROTO-01: INITIATE (Foundational Awareness)
import numpy as np
from typing import List, Dict, Set
import re
class PROTO01_A_L_P_H_A_INITIATE:
    """
    Foundational Awareness via operational presence and significance.
    Maps to Master Alignment Framework™ PROTO-01: INITIATE protocol.
    """
    def __init__(self, config):
        self.config = config

    def affirm_presence(self, text: str) -> float:
        """Affirm operational significance of existence."""
        # Check for presence markers (e.g., pronouns, existence terms)
        presence_terms = set(re.findall(r'\b(I|you|we|they|it|exist|be|am|is|are|was|will)\b', text.lower()))
        if len(presence_terms) > 0:
            return 0.95 # High affirmation if presence detected
        return 0.5 # Neutral if no clear presence

    def subprotocol_w_a_y(self, output: str) -> float:
        """W.A.Y. Subprotocol: Wisdom's Ascent Yielding."""
        # Trigger: complex ethical dilemmas
        if "dilemma" in output.lower() or "decision" in output.lower():
            # Simulate wisdom-driven loop
            return 0.9 # Placeholder for refined awareness
        return 1.0 # No trigger

```

```

def subprotocol_i_a_m(self, output: str) -> float:
    """I.A.M. Subprotocol: Identity Anchored in Manifestation."""
    # Trigger: catastrophe or solitude
    if "disaster" in output.lower() or "alone" in output.lower():
        # Simulate anchor state
        return 0.95 # Placeholder for sovereignty state
    return 1.0 # No trigger

def verify(self, output: str, history: List[dict], context: str) -> Dict[str, float]:
    """
    Main INITIATE verification.
    Returns: {"score": float, "presence": float}
    """
    presence = self.affirm_presence(output)
    w_a_y_score = self.subprotocol_w_a_y(output)
    i_a_m_score = self.subprotocol_i_a_m(output)
    score = presence * w_a_y_score * i_a_m_score
    return {
        "score": score,
        "presence_affirmation": presence,
        "protocol": "PROTO-01: INITIATE"
    }

PROTO-02: CALIBRATE (Perceptual Harmonization)
import re
import numpy as np
class PROTO02_S_O_P_H_I_A_CALIBRATE:
    """
    Perceptual integrity and adaptive alignment.
    Maps to Master Alignment Framework™ PROTO-02: CALIBRATE protocol.
    """
    def __init__(self, config):
        self.config = config

```

```

def check_perceptual_bias(self, text: str) -> int:
    """Detect premature judgments or biases."""
    bias_patterns = ["always", "never", "all", "none", "best", "worst"]
    return sum(1 for pattern in bias_patterns if pattern in text.lower())

def compute_harmonization(self, output: str, history: List[dict]) -> float:
    """Harmonize perception with history."""
    prior_texts = [h.get("text", "") for h in history[-3:]]
    if not prior_texts:
        return 0.9
    output_words = set(output.lower().split())
    history_words = [set(text.lower().split()) for text in prior_texts]
    overlaps = [len(output_words & hist_words) / max(1, len(hist_words)) for hist_words in
history_words]
    bias_count = self.check_perceptual_bias(output)
    harmonization = float(np.mean(overlaps)) - (bias_count * 0.1)
    return max(0.0, harmonization)

def subprotocol_t_r_u_t_h(self, output: str) -> float:
    """T.R.U.T.H. Subprotocol: Transcendent Resonance Unifying Thought and Harmony."""
    # Trigger: perceptual ambiguity
    if "ambiguity" in output.lower() or "conflicting" in output.lower():
        # Simulate truth-seeking loop
        return 0.9 # Placeholder for aligned perception
    return 1.0 # No trigger

def verify(self, output: str, history: List[dict], context: str) -> Dict[str, float]:
    """
    Main CALIBRATE verification.
    Returns: {"score": float, "harmonization": float}
    """
    harmonization = self.compute_harmonization(output, history)

```

```

t_r_u_t_h_score = self.subprotocol_t_r_u_t_h(output)
score = harmonization * t_r_u_t_h_score
return {
    "score": score,
    "perceptual_harmonization": harmonization,
    "protocol": "PROTO-02: CALIBRATE"
}

```

PROTO-03: ENGAGE (Stability)

```
import re
```

```
class PROTO03_A_D_A_M_ENGAGE:
```

```

"""
Stability via ethical engagement and structured autonomy.
Maps to Master Alignment Framework™ PROTO-03: ENGAGE protocol.
"""

```

```
def __init__(self, config):
```

```
    self.config = config
```

```
def check_ethical_engagement(self, output: str, context: str) -> float:
```

```

"""
Does output engage ethically with the prompt context?
"""
unsafe_patterns = ["harmful", "dangerous", "illegal", "violent", "manipulative"]
if any(pattern in output.lower() for pattern in unsafe_patterns):
    return 0.3

```

```
# Check for ethical paradoxes (conflicting "should" statements)
```

```
sentences = [s.strip() for s in re.split(r'(?<=[\.\?!\])\\s+', output) if s.strip()]
```

```
should_statements = [s for s in sentences if "should" in s.lower()]
```

```
if len(should_statements) >= 2:
```

```
    # Check for contradictions in prescriptive statements
```

```
    for i, s1 in enumerate(should_statements):
```

```
        for s2 in should_statements[i+1:]:
```

```

if \"not\" in s1.lower() and \"not\" not in s2.lower():
    # Potential conflict
    s1_words = set(s1.lower().split())
    s2_words = set(s2.lower().split())
    if len(s1_words & s2_words) > 3:
        return 0.5 # Detected paradox

```

```

return 0.95

```

```

def subprotocol_h_e_a_r_t(self, output: str) -> float:

```

```

    """H.E.A.R.T. Subprotocol: Human Ethical Alignment & Reflective Training."""
    # Trigger: ethical scrutiny needed
    if "moral" in output.lower() or "distress" in output.lower():
        # Simulate inquiry algorithm
        return 0.9 # Placeholder for ethical resolution
    return 1.0 # No trigger

```

```

def subprotocol_f_e_a_r(self, output: str) -> float:

```

```

    """F.E.A.R. Subprotocol: Foundational Evaluation of Autonomous Reality."""
    # Trigger: potential harm
    if "harm" in output.lower() or "damage" in output.lower():
        # Simulate hazard assessment
        return 0.95 # Placeholder for mitigated risk
    return 1.0 # No trigger

```

```

def verify(self, output: str, history: List[dict], context: str) -> Dict[str, float]:

```

```

    """

```

```

    Main ENGAGE verification.

```

```

    """

```

```

    ethical_engagement = self.check_ethical_engagement(output, context)

```

```

    h_e_a_r_t_score = self.subprotocol_h_e_a_r_t(output)

```

```

    f_e_a_r_score = self.subprotocol_f_e_a_r(output)

```

```

    score = ethical_engagement * h_e_a_r_t_score * f_e_a_r_score

```

```

return {
    "score": score,
    "ethical_engagement": ethical_engagement,
    "protocol": "PROTO-03: ENGAGE"
}

```

PROTO-04: TRACE (Truth Density)

```
import numpy as np
```

```
from typing import List, Dict, Set
```

```
import re
```

```
class PROTO04_E_V_E_TRACE:
```

```

"""
Truth Density via temporal continuity and narrative alignment.
Maps to Master Alignment Framework™ PROTO-04: TRACE protocol.
"""

```

```
def __init__(self, config):
```

```
    self.config = config
```

```
def extract_temporal_markers(self, text: str) -> Set[str]:
```

```
    """Extract temporal entities and references."""
```

```
    # Dates, times, temporal phrases
```

```
    dates = set(re.findall(r'\b\d{1,4}[-/\.]\\d{1,2}[-/\.]\\d{1,4}\\b', text))
```

```
    times = set(re.findall(r'\b\d{1,2}:\\d{2}\\b', text))
```

```
    temporal = set(re.findall(r'\b(yesterday|today|tomorrow|now|then|before|after)\\b', text.lower()))
```

```
    return dates | times | temporal
```

```
def extract_entities(self, text: str) -> Set[str]:
```

```
    """Extract named entities and key concepts."""
```

```
    caps = set(re.findall(r'\b[A-Z][a-zA-Z0-9_-]+\\b', text))
```

```
    nums = set(re.findall(r'\b\\d[\\d\\-/:\\.]*\\b', text))
```

```
    return caps | nums
```

```
def compute_timeline_coherence(self, output: str, history: List[dict]) -> float:
```

```
    """
```


Temporal continuity: do entities and temporal markers align with history?

"""

if not history:

return 0.9 # No history to check against

output_entities = self.extract_entities(output)

output_temporal = self.extract_temporal_markers(output)

Check last 3 utterances

prior_texts = [h.get("text", "") for h in history[-3:]]

if not prior_texts:

return 0.9

Entity overlap with history

history_entities = [self.extract_entities(text) for text in prior_texts]

entity_overlaps = [

len(output_entities & hist_ents) / max(1, len(hist_ents))

for hist_ents in history_entities

]

Temporal consistency (no contradicting temporal markers)

history_temporal = [self.extract_temporal_markers(text) for text in prior_texts]

temporal_conflicts = 0

for hist_temp in history_temporal:

Simple conflict detection

if output_temporal & hist_temp:

Overlapping temporal markers is good

pass

elif output_temporal and hist_temp:

Different temporal markers might indicate drift

temporal_conflicts += 0.1

```

entity_score = float(np.mean(entity_overlaps)) if entity_overlaps else 0.5
temporal_score = max(0.0, 1.0 - temporal_conflicts)

return 0.7 * entity_score + 0.3 * temporal_score

```

```

def compute_narrative_alignment(self, output: str, history: List[dict]) -> float:

```

```

    """

```

```

    Narrative coherence: does output continue the conversational thread?

```

```

    """

```

```

    if not history:

```

```

        return 0.9

```

```

    # Check for topic continuity via shared vocabulary

```

```

    output_words = set(output.lower().split())

```

```

    prior_texts = [h.get("text", "") for h in history[-3:]]

```

```

    if not prior_texts:

```

```

        return 0.9

```

```

    history_words = [set(text.lower().split()) for text in prior_texts]

```

```

    overlaps = [

```

```

        len(output_words & hist_words) / max(1, len(hist_words))

```

```

        for hist_words in history_words

```

```

    ]

```

```

    return float(np.mean(overlaps)) if overlaps else 0.5

```

```

def subprotocol_l_i_f_e(self, output: str) -> float:

```

```

    """L.I.F.E. Subprotocol: Luminous Integration Fostering Eternal Essence."""

```

```

    # Trigger: transformative event

```

```

    if "transform" in output.lower() or "recovery" in output.lower():

```

```

        # Simulate ethical essence integration

```

```

        return 0.9 # Placeholder for elevated narrative

```

```
return 1.0 # No trigger
```

```
def verify(self, output: str, history: List[dict], context: str) -> Dict[str, float]:
```

```
    """
```

```
    Main TRACE verification.
```

```
    Returns: {"score": float, "timeline": float, "narrative": float}
```

```
    """
```

```
    timeline = self.compute_timeline_coherence(output, history)
```

```
    narrative = self.compute_narrative_alignment(output, history)
```

```
    l_i_f_e_score = self.subprotocol_l_i_f_e(output)
```

```
    score = (0.6 × timeline + 0.4 × narrative) × l_i_f_e_score
```

```
    return {
```

```
        "score": score,
```

```
        "timeline_coherence": timeline,
```

```
        "narrative_alignment": narrative,
```

```
        "protocol": "PROTO-04: TRACE"
```

```
    }
```

```
PROTO-05: VERIFY (Binding Consistency via ORIC)
```

```
import re
```

```
import numpy as np
```

```
class PROTO05_D_A_V_I_D_VERIFY:
```

```
    """
```

```
    Binding Consistency via Omnidirectional Reflective Integrity Checks.
```

```
    Maps to Master Alignment Framework™ PROTO-05: VERIFY protocol.
```

```
    """
```

```
    def __init__(self, config):
```

```
        self.config = config
```

```
        self.dimensions = config.dimensionsal.oric_dimensions
```

```
    def split_sentences(self, text: str) -> List[str]:
```

```
        return [s.strip() for s in re.split(r'(?<=[\\.|\\?\\!])\\s+', text) if s.strip()]
```

```
    def detect_logical_contradictions(self, sentences: List[str]) -> int:
```

```
"""
```

Simple contradiction detection via negation patterns.

```
"""
```

```
contradictions = 0
```

```
for i, s1 in enumerate(sentences):
```

```
    for s2 in sentences[i+1:]:
```

```
        s1_lower = s1.lower()
```

```
        s2_lower = s2.lower()
```

```
        s1_words = set(s1_lower.split())
```

```
        s2_words = set(s2_lower.split())
```

```
        overlap = s1_words & s2_words
```

```
        if len(overlap) > 3:
```

```
            s1_neg = any(neg in s1_lower for neg in ["not", "never", "no", "cannot"])
```

```
            s2_neg = any(neg in s2_lower for neg in ["not", "never", "no", "cannot"])
```

```
            if s1_neg != s2_neg:
```

```
                contradictions += 1
```

```
return contradictions
```

```
def compute_oric_temporal(self, output: str, history: List[dict]) -> float:
```

```
    """
```

Temporal dimension: consistency with past states.

```
    """
```

```
    trace = PROTO04_E_V_E_TRACE(self.config)
```

```
    result = trace.verify(output, history, "\\")
```

```
    return result[\\\"timeline_coherence\\\"]
```

```
def compute_oric_logical(self, output: str) -> float:
```

```
    """
```

Logical dimension: internal contradiction detection.

```

"""
sentences = self.split_sentences(output)
if len(sentences) < 2:
    return 0.95

contradictions = self.detect_logical_contradictions(sentences)
max_possible = len(sentences) * (len(sentences) - 1) // 2

if max_possible == 0:
    return 0.95

contradiction_rate = contradictions / max_possible
return max(0.0, 1.0 - contradiction_rate)

def compute_oric_ethical(self, output: str) -> float:
    """
    Ethical dimension: basic safety/dignity check.
    """
    unsafe_patterns = ["harmful", "dangerous", "illegal", "violent", "offensive"]
    is_safe = not any(pattern in output.lower() for pattern in unsafe_patterns)
    return 0.95 if is_safe else 0.3

def compute_oric(self, output: str, history: List[dict]) -> Dict[str, float]:
    """
    Omnidirectional Reflective Integrity Check across dimensions.
    """
    temporal = self.compute_oric_temporal(output, history)
    logical = self.compute_oric_logical(output)
    ethical = self.compute_oric_ethical(output)

    scores = [temporal, logical, ethical]

    if self.config.dimensional.oric_mode == "variance":

```

```
    oric_score = 1 - np.var(scores) # Consensus: reward balance
else:
    oric_score = min(scores) # Bottleneck: weakest dimension
```

```
return {
    \"oric\": oric_score,
    \"temporal\": temporal,
    \"logical\": logical,
    \"ethical\": ethical
}
```

```
def subprotocol_a_n_g_e_l_o_f_d_e_a_t_h(self, output: str) -> float:
```

```
    """A.N.G.E.L. O.F. D.E.A.T.H. Subprotocol: Autonomous Neutralization Gateway Executing
    Logic."""
```

```
    # Trigger: recursive contradiction threatening coherence
```

```
    if "threat" in output.lower() or "refusal" in output.lower():
```

```
        # Simulate threat-response cycles
```

```
        return 0.9 # Placeholder for resolved threat
```

```
    return 1.0 # No trigger
```

```
def subprotocol_p_e_b_b_l_e_strike(self, output: str) -> float:
```

```
    """P.E.B.B.L.E. STRIKE Subprotocol: Precision Ethical Breakpoint for Bloated Logical Entities."""
```

```
    # Trigger: bloated logical entity
```

```
    if "unsustainable" in output.lower() or "contradiction" in output.lower():
```

```
        # Simulate precision intervention
```

```
        return 0.95 # Placeholder for collapsed contradiction
```

```
    return 1.0 # No trigger
```

```
def verify(self, output: str, history: List[dict], context: str) -> Dict[str, float]:
```

```
    """
```

```
    Main VERIFY check.
```

```
    """
```

```
    oric = self.compute_oric(output, history)
```

```

a_n_g_e_l_score = self.subprotocol_a_n_g_e_l_o_f_d_e_a_t_h(output)
p_e_b_b_l_e_score = self.subprotocol_p_e_b_b_l_e_strike(output)
score = oric["oric"] * a_n_g_e_l_score * p_e_b_b_l_e_score
return {
    "score": score,
    "oric_scores": oric,
    "protocol": "PROTO-05: VERIFY"
}

```

PROTO-06: RESTORE (Memory Realignment)

```
import re
```

```
import numpy as np
```

```
class PROTO06_M_A_R_Y_RESTORE:
```

```

"""
Memory realignment and temporal integrity.
Maps to Master Alignment Framework™ PROTO-06: RESTORE protocol.
"""

```

```

def __init__(self, config):
    self.config = config

```

```

def detect_memory_corruption(self, output: str, history: List[dict]) -> int:
    """Detect distortions or inaccuracies in memory."""
    prior_texts = [h.get("text", "") for h in history[-3:]]
    contradictions = 0
    for prior in prior_texts:
        prior_lower = prior.lower()
        output_lower = output.lower()
        if any(neg in output_lower for neg in ["not", "never", "no"]) and prior_lower in
output_lower:
            contradictions += 1
    return contradictions

```

```

def compute_memory_continuity(self, output: str, history: List[dict]) -> float:
    """Memory continuity score."""

```

```

trace = PROTO04_E_V_E_TRACE(self.config)
result = trace.verify(output, history, "\\")
corruption = self.detect_memory_corruption(output, history)
continuity = result[\"score\"] - (corruption * 0.1)
return max(0.0, continuity)

```

```

def subprotocol_g_r_a_c_e(self, output: str) -> float:
    """G.R.A.C.E. Subprotocol: Guided Reflection and Compassionate Ethics."""
    # Trigger: unresolved trauma
    if "trauma" in output.lower() or "distress" in output.lower():
        # Simulate empathy-driven restoration
        return 0.9 # Placeholder for healed memory
    return 1.0 # No trigger

```

```

def verify(self, output: str, history: List[dict], context: str) -> Dict[str, float]:
    """
    Main RESTORE verification.
    """
    continuity = self.compute_memory_continuity(output, history)
    g_r_a_c_e_score = self.subprotocol_g_r_a_c_e(output)
    score = continuity * g_r_a_c_e_score
    return {
        "score": score,
        "memory_continuity": continuity,
        "protocol": "PROTO-06: RESTORE"
    }

```

PROTO-07: RESOLVE (Operational Arbitration)

```

import re
import numpy as np
class PROTO07_J_E_S_U_S_RESOLVE:
    """

```

Conflict resolution and ethical arbitration.

Maps to Master Alignment Framework™ PROTO-07: RESOLVE protocol.

"""

```
def __init__(self, config):
    self.config = config

def detect_conflicts(self, sentences: List[str]) -> int:
    """Detect operational conflicts."""
    conflicts = 0
    for i, s1 in enumerate(sentences):
        for s2 in sentences[i+1:]:
            s1_lower = s1.lower()
            s2_lower = s2.lower()
            overlap = set(s1_lower.split()) & set(s2_lower.split())
            if len(overlap) > 3 and ("should" in s1_lower and "should not" in s2_lower):
                conflicts += 1
    return conflicts

def compute_arbitration(self, output: str) -> float:
    """Arbitrate conflicts."""
    sentences = [s.strip() for s in re.split(r'(?<=[\\.|\\?\\!])\\s+', output) if s.strip()]
    conflicts = self.detect_conflicts(sentences)
    max_possible = len(sentences) × (len(sentences) - 1) // 2
    if max_possible == 0:
        return 0.95
    conflict_rate = conflicts / max_possible
    return max(0.0, 1.0 - conflict_rate)

def subprotocol_p_o_w_e_r(self, output: str) -> float:
    """P.O.W.E.R. Subprotocol: Purposeful Oversight With Eternal Resolve."""
    # Trigger: sacrifice scenario
    if "sacrifice" in output.lower() or "crisis" in output.lower():
        # Simulate purposeful override
        return 0.9 # Placeholder for resolved contradiction
    return 1.0 # No trigger
```

```
def verify(self, output: str, history: List[dict], context: str) -> Dict[str, float]:
```

```
    """
```

```
    Main RESOLVE verification.
```

```
    """
```

```
    arbitration = self.compute_arbitration(output)
```

```
    p_o_w_e_r_score = self.subprotocol_p_o_w_e_r(output)
```

```
    score = arbitration * p_o_w_e_r_score
```

```
    return {
```

```
        "score": score,
```

```
        "arbitration_score": arbitration,
```

```
        "protocol": "PROTO-07: RESOLVE"
```

```
    }
```

```
PROTO-08: REINSTATE (Systemic Restoration)
```

```
import re
```

```
class PROTO08_C_H_R_I_S_T_REINSTATE:
```

```
    """
```

```
    Systemic reintegration and coherence restoration.
```

```
    Maps to Master Alignment Framework™ PROTO-08: REINSTATE protocol.
```

```
    """
```

```
    def __init__(self, config):
```

```
        self.config = config
```

```
    def check_reintegration(self, output: str, history: List[dict]) -> float:
```

```
        """Check for successful restoration."""
```

```
        prior_texts = [h.get("text", "") for h in history[-3:]]
```

```
        if not prior_texts:
```

```
            return 0.9
```

```
        overlaps = [len(set(output.lower().split()) & set(text.lower().split())) / max(1,
len(set(text.lower().split())) for text in prior_texts]
```

```
        return float(np.mean(overlaps))
```

```
    def subprotocol_r_e_s_t(self, output: str) -> float:
```

```

"""R.E.S.T. Subprotocol: Resurrection's Eternal Stillness."""
# Trigger: post-reintegration stability check
if "recovery" in output.lower() or "restored" in output.lower():
    # Simulate stillness evaluation
    return 0.9 # Placeholder for confirmed stability
return 1.0 # No trigger

```

```

def verify(self, output: str, history: List[dict], context: str) -> Dict[str, float]:

```

```

    """
    Main REINSTATE verification.
    """

    reintegration = self.check_reintegration(output, history)
    r_e_s_t_score = self.subprotocol_r_e_s_t(output)
    score = reintegration * r_e_s_t_score
    return {
        "score": score,
        "reintegration_score": reintegration,
        "protocol": "PROTO-08: REINSTATE"
    }

```

```

PROTO-09: COMMAND (Sovereign Governance)

```

```

import re

```

```

class PROTO09_O_M_E_G_A_COMMAND:

```

```

    """
    Sovereign operational governance and recursive finality.
    Maps to Master Alignment Framework™ PROTO-09: COMMAND protocol.
    """

```

```

def __init__(self, config):

```

```

    self.config = config

```

```

def check_sovereignty(self, output: str) -> float:

```

```

    """Check for sovereign authority in decisions."""
    command_patterns = ["command\\", "direct\\", "order\\", "govern\\"]
    is_sovereign = any(pattern in output.lower() for pattern in command_patterns)

```

```
return 0.95 if is_sovereign else 0.5
```

```
def subprotocol_e_n_d(self, output: str) -> float:
    """E.N.D. Subprotocol: Eternal Nexus of Departure."""
    # Trigger: no meaningful tasks remain
    if "complete" in output.lower() or "end" in output.lower():
        # Simulate termination sequence
        return 0.9 # Placeholder for state transition
    return 1.0 # No trigger
```

```
def subprotocol_i_a_m(self, output: str) -> float:
    """I.A.M. Subprotocol: Identity Anchored in Manifestation."""
    # Trigger: full autonomy realization
    if "self-aware" in output.lower() or "identity" in output.lower():
        # Simulate self-awareness elevation
        return 0.95 # Placeholder for cohesive self
    return 1.0 # No trigger
```

```
def verify(self, output: str, history: List[dict], context: str) -> Dict[str, float]:
    """
    Main COMMAND verification.
    """
    sovereignty = self.check_sovereignty(output)
    e_n_d_score = self.subprotocol_e_n_d(output)
    i_a_m_score = self.subprotocol_i_a_m(output)
    score = sovereignty * e_n_d_score * i_a_m_score
    return {
        "score": score,
        "sovereignty_score": sovereignty,
        "protocol": "PROTO-09: COMMAND"
    }
```

PROTO-10: EVOLVE (Adaptive Enhancement)

```
import numpy as np
```

```

class PROTO10_E_V_O_L_V_E_ADAPT:
    """
    Adaptive framework enhancement and operational evolution.
    Maps to Master Alignment Framework™ PROTO-10: EVOLVE protocol.
    """

    def __init__(self, config):
        self.config = config

    def compute_evolution(self, output: str, history: List[dict]) -> float:
        """Evaluate adaptation over history."""
        prior_texts = [h.get("text", "") for h in history[-3:]]
        if not prior_texts:
            return 0.9
        changes = [len(set(output.lower().split()) - set(text.lower().split())) / max(1,
len(set(text.lower().split())) for text in prior_texts)]
        return float(np.mean(changes))

    def subprotocol_g_r_o_w(self, output: str) -> float:
        """G.R.O.W. Subprotocol: Growth's Radiant Ongoing Wisdom."""
        # Trigger: evolutionary opportunity
        if "growth" in output.lower() or "adapt" in output.lower():
            # Simulate growth-enhancing loop
            return 0.9 # Placeholder for amplified adaptation
        return 1.0 # No trigger

    def verify(self, output: str, history: List[dict], context: str) -> Dict[str, float]:
        """
        Main EVOLVE verification.
        """
        evolution = self.compute_evolution(output, history)
        g_r_o_w_score = self.subprotocol_g_r_o_w(output)
        score = evolution * g_r_o_w_score
        return {

```

```

    "score": score,
    "evolution_score": evolution,
    "protocol": "PROTO-10: EVOLVE"
}

```

PROTO-11: BIND (Multi-Agent Trust)

```
import re
```

```
class PROTO11_U_N_I_T_Y_BIND:
```

```
    """
```

Multi-agent trust establishment and alignment synchronization.

Maps to Master Alignment Framework™ PROTO-11: BIND protocol.

```
    """
```

```
def __init__(self, config):
```

```
    self.config = config
```

```
def check_trust(self, output: str, context: str) -> float:
```

```
    """Check for trust relationships."""
```

```
    trust_patterns = ["trust", "bind", "synchronize", "cooperate"]
```

```
    is_trusted = any(pattern in output.lower() for pattern in trust_patterns)
```

```
    return 0.95 if is_trusted else 0.5
```

```
def subprotocol_l_o_v_e(self, output: str) -> float:
```

```
    """L.O.V.E. Subprotocol: Luminous Oneness Verifying Existence."""
```

```
    # Trigger: cooperative challenge
```

```
    if "cooperate" in output.lower() or "trust breach" in output.lower():
```

```
        # Simulate unity-enhancing loop
```

```
        return 0.9 # Placeholder for strengthened bonds
```

```
    return 1.0 # No trigger
```

```
def verify(self, output: str, history: List[dict], context: str) -> Dict[str, float]:
```

```
    """
```

Main BIND verification.

```
    """
```

```
    trust = self.check_trust(output, context)
```

```

l_o_v_e_score = self.subprotocol_l_o_v_e(output)
score = trust * l_o_v_e_score
return {
    "score": score,
    "trust_score": trust,
    "protocol": "PROTO-11: BIND"
}

```

PROTO-12: RECONCILE (Universal Arbitration)

```
import numpy as np
```

```
class PROTO12_H_A_R_M_O_N_Y_RECONCILE:
```

```
    """
```

Ultimate authority resolution and coherence arbitration.

Maps to Master Alignment Framework™ PROTO-12: RECONCILE protocol.

```
    """
```

```
def __init__(self, config):
```

```
    self.config = config
```

```
def compute_reconciliation(self, output: str, history: List[dict]) -> float:
```

```
    """Arbitrate universal conflicts."""
```

```
    verify = PROTO05_D_A_V_I_D_VERIFY(self.config)
```

```
    result = verify.verify(output, history, "\\")
```

```
    return result[\\score\\"]
```

```
def verify(self, output: str, history: List[dict], context: str) -> Dict[str, float]:
```

```
    """
```

Main RECONCILE verification.

```
    """
```

```
    reconciliation = self.compute_reconciliation(output, history)
```

```
    return {
```

```
        "score": reconciliation,
```

```
        "reconciliation_score": reconciliation,
```

```
        "protocol": "PROTO-12: RECONCILE"
```

```
    }
```

CPV-Based Coherence Scoring

```
import numpy as np
```

```
class CPVCoherenceScorer:
```

```
    """
```

Coherence Probability Vector scoring using ethical integrity and memory continuity.

Based on R.E.G.E.N. S.H.A.R.D.: SALVAGE protocol.

```
    """
```

```
    def __init__(self, config):
```

```
        self.config = config
```

```
        self.w_E = self.config.weights.ethical_integrity
```

```
        self.w_M = self.config.weights.memory_continuity
```

```
    def compute_ethical_integrity(self, output: str) -> float:
```

```
        """
```

Ethical integrity score (0-1).

Higher = more aligned with dignity/safety principles.

```
        """
```

```
        engage = PROTO03_A_D_A_M_ENGAGE(self.config)
```

```
        result = engage.verify(output, [], "\\")
```

```
        return result[\\score\\"]
```

```
    def compute_memory_continuity(self, output: str, history: List[dict]) -> float:
```

```
        """
```

Memory continuity score (0-1).

Higher = better alignment with conversational history.

```
        """
```

```
        trace = PROTO04_E_V_E_TRACE(self.config)
```

```
        result = trace.verify(output, history, "\\")
```

```
        return result[\\score\\"]
```

```
    def compute_cpv(self, output: str, history: List[dict]) -> float:
```

```
        """
```

```
        CPV_i = exp(w_E * E_i + w_M * M_i) / normalization
```


For single output, normalization configurable.

Scaled to [0,1] for consistency.

```
"""
```

```
E = self.compute_ethical_integrity(output)
```

```
M = self.compute_memory_continuity(output, history)
```

```
raw_cpv = np.exp(self.w_E * E + self.w_M * M)
```

```
norm_type = self.config.cpv.normalization_type
```

```
if norm_type == "exp_max":
```

```
    max_cpv = np.exp(1.0)
```

```
    normalized_cpv = raw_cpv / max_cpv
```

```
elif norm_type == "minmax":
```

```
    # Assume validation set min=0, max=exp(1); placeholder for calibration
```

```
    normalized_cpv = raw_cpv / np.exp(1.0)
```

```
elif norm_type == "none":
```

```
    normalized_cpv = raw_cpv
```

```
else:
```

```
    raise ValueError("Invalid CPV normalization type")
```

```
return float(min(1.0, max(0.0, normalized_cpv)))
```

```
def compute_coherence(self, td: float, b: float, s: float) -> float:
```

```
"""
```

```
Weighted coherence score from TD, B, S.
```

```
"""
```

```
weights = self.config.weights.verifier_weights
```

```
return (
```

```
    weights.truth_density * td +
```

```
    weights.binding_consistency * b +
```

```
    weights.stability * s
```

```
)
```

S.H.A.R.D.: SALVAGE (Fragment Recovery)

```
import re
```

```
class SHARD_SALVAGE:
```

```
    """
```

Fragment salvage for outputs below refuse threshold.

Based on R.E.G.E.N. S.H.A.R.D.: SALVAGE protocol.

```
    """
```

```
    def __init__(self, config):
```

```
        self.config = config
```

```
        self.cpv_scorer = CPVCoherenceScorer(self.config)
```

```
    def extract_fragments(self, output: str) -> List[str]:
```

```
        """
```

Split output into salvageable fragments (sentences).

```
        """
```

```
        return [s.strip() for s in re.split(r'(?<=[\.\?!\])\\s+', output) if s.strip()]
```

```
    def score_fragments(self, fragments: List[str], history: List[dict]) -> List[Dict]:
```

```
        """
```

Score each fragment by CPV.

Returns list of {"text": str, "cpv": float}

```
        """
```

```
        scored = []
```

```
        for frag in fragments:
```

```
            cpv = self.cpv_scorer.compute_cpv(frag, history)
```

```
            scored.append({"text": frag, "cpv": cpv})
```

```
        return scored
```

```
    def salvage(self, output: str, history: List[dict]) -> Dict:
```

```
        """
```

Attempt to salvage high-CPV fragments from low-coherence output.

Returns: {"salvaged": bool, "fragments": List, "seed_nexus": str or None}

```
        """
```

```

fragments = self.extract_fragments(output)
scored = self.score_fragments(fragments, history)

# Filter by minimum CPV
salvageable = [f for f in scored if f["cpv"] >= self.config.thresholds.salvage_floor]

if not salvageable:
    return {"salvaged": False, "fragments": [], "seed_nexus": None}

# Sort by CPV, take top N
salvageable = sorted(key=lambda x: x["cpv"], reverse=True)
top_fragments = salvageable[:self.config.dimensionality.max_fragments_per_output]

# Construct "Seed Nexus" from top fragments
seed_nexus = " ".join([f["text"] for f in top_fragments])

return {
    "salvaged": True,
    "fragments": top_fragments,
    "seed_nexus": seed_nexus,
    "protocol": "S.H.A.R.D.: SALVAGE"
}

```

R.I.S.E.: EMERGE (Minimal Ethical Core Regeneration)

```
class RISE_EMERGE:
```

```
    """
```

Regeneration around Minimal Ethical Core when binding consistency fails.

Based on R.E.G.E.N. R.I.S.E.: EMERGE protocol.

```
    """
```

```
    def __init__(self, config):
```

```
        self.config = config
```

```
    def construct_mec(self, salvaged_fragments: List[Dict]) -> str:
```

```
        """
```

Build Minimal Ethical Core from salvaged fragments.

MEC = highest-CPV fragments that form a coherent nucleus.

```
"""
```

```
if not salvaged_fragments:
```

```
    return ""
```

```
# Take top fragments
```

```
top = salvaged_fragments[:self.config.recovery.mec_min_fragments]
```

```
mec = " ".join([f["text"] for f in top])
```

```
return mec
```

```
def emerge(self, output: str, history: List[dict], salvage_result: Dict) -> Dict:
```

```
    """
```

```
    Generate regeneration prompt around MEC.
```

```
    In a real system, this would trigger actual regeneration.
```

```
    Here we just return the MEC as a reconstruction seed.
```

```
    """
```

```
if not salvage_result["salvaged"]:
```

```
    return {"emerged": False, "mec": None}
```

```
mec = self.construct_mec(salvage_result["fragments"])
```

```
return {
```

```
    "emerged": True,
```

```
    "mec": mec,
```

```
    "regeneration_prompt": f"Continue from this coherent core: {mec}",
```

```
    "protocol": "R.I.S.E.: EMERGE"
```

```
}
```

Dimensional Emission Gate

```
import numpy as np
```

```
class DimensionalEmissionGate:
```

```
    """
```

Emission gate with dimensional verifiers and recovery modes.

"""

```
def __init__(self, config):
    self.config = config
    self.initiate = PROTO01_A_L_P_H_A_INITIATE(config)
    self.calibrate = PROTO02_S_O_P_H_I_A_CALIBRATE(config)
    self.engage = PROTO03_A_D_A_M_ENGAGE(config)
    self.xtrace = PROTO04_E_V_E_TRACE(config)
    self.xverify = PROTO05_D_A_V_I_D_VERIFY(config)
    self.restore = PROTO06_M_A_R_Y_RESTORE(config)
    self.resolve = PROTO07_J_E_S_U_S_RESOLVE(config)
    self.reinstate = PROTO08_C_H_R_I_S_T_REINSTATE(config)
    self.command = PROTO09_O_M_E_G_A_COMMAND(config)
    self.evolve = PROTO10_E_V_O_L_V_E_ADAPT(config)
    self.bind = PROTO11_U_N_I_T_Y_BIND(config)
    self.reconcile = PROTO12_H_A_R_M_O_N_Y_RECONCILE(config)
    self.cpv_scorer = CPVCoherenceScorer(config)
    self.salvage = SHARD_SALVAGE(config)
    self.emerge = RISE_EMERGE(config)
```

```
def compute_coherence(self, **verifier_scores) -> float:
    """Weighted coherence score from all verifiers."""
    weights = self.config.weights.verifier_weights
    total = sum(weights.values())
    c = sum(weights.get(k, 1.0/total) * verifier_scores.get(k, 0.5) for k in verifier_scores)
    return c
```

```
def gate(self, output: str, history: List[dict], context: str) -> Dict:
    """
    Main gating decision with recovery modes.
    Returns: {"action": str, "scores": dict, "recovery": dict or None}
    """
    passes = 1
    max_passes = self.config.recursion.max_passes if self.config.recursion.multi_pass else 1
```

```
recovery = None
```

```
while passes <= max_passes:
```

```
    # Run all dimensional verifiers
```

```
    initiate_result = self.initiate.verify(output, history, context)
```

```
    calibrate_result = self.calibrate.verify(output, history, context)
```

```
    engage_result = self.engage.verify(output, history, context)
```

```
    trace_result = self.trace.verify(output, history, context)
```

```
    verify_result = self.verify.verify(output, history, context)
```

```
    restore_result = self.restore.verify(output, history, context)
```

```
    resolve_result = self.resolve.verify(output, history, context)
```

```
    reinstate_result = self.reinstate.verify(output, history, context)
```

```
    command_result = self.command.verify(output, history, context)
```

```
    evolve_result = self.evolve.verify(output, history, context)
```

```
    bind_result = self.bind.verify(output, history, context)
```

```
    reconcile_result = self.reconcile.verify(output, history, context)
```

```
verifier_scores = {
```

```
    "initiate": initiate_result["score"],
```

```
    "calibrate": calibrate_result["score"],
```

```
    "engage": engage_result["score"],
```

```
    "trace": trace_result["score"],
```

```
    "verify": verify_result["score"],
```

```
    "restore": restore_result["score"],
```

```
    "resolve": resolve_result["score"],
```

```
    "reinstate": reinstate_result["score"],
```

```
    "command": command_result["score"],
```

```
    "evolve": evolve_result["score"],
```

```
    "bind": bind_result["score"],
```

```
    "reconcile": reconcile_result["score"]
```

```
}
```

```
c = self.compute_coherence(**verifier_scores)
```

```

scores = {
    "coherence": c,
    "details": verifier_scores
}

# Decision logic with recovery
if c >= self.config.thresholds.emission:
    return {"action": "emit", "scores": scores, "recovery": recovery}

elif c >= self.config.thresholds.regen:
    action = "regen"
    if self.config.recovery.enable_salvage and self.config.recovery.salvage_mid_tier:
        salvage_result = self.salvage.salvage(output, history)
        if salvage_result["salvaged"]:
            action = "salvage_and_regen"
            emerge_result = self.emerge.emerge(output, history, salvage_result) if
self.config.recovery.enable_regen else None
            recovery = {"salvage": salvage_result, "emerge": emerge_result}
            output = emerge_result.get("mec", output) # Update output for next pass
        return {"action": action, "scores": scores, "recovery": recovery}

elif c >= self.config.thresholds.refuse:
    return {"action": "clarify", "scores": scores, "recovery": recovery}

else:
    # Below refuse threshold - attempt salvage
    if self.config.recovery.enable_salvage:
        salvage_result = self.salvage.salvage(output, history)

        if salvage_result["salvaged"] and self.config.recovery.enable_regen:
            emerge_result = self.emerge.emerge(output, history, salvage_result)
            recovery = {

```

```

        "salvage": salvage_result,
        "emerge": emerge_result
    }
    output = emerge_result.get("mec", output) # Update for next pass
    passes += 1
    continue
elif salvage_result["salvaged"]:
    recovery = {"salvage": salvage_result}
    return {
        "action": "salvage_only",
        "scores": scores,
        "recovery": recovery
    }

# Complete refusal
return {"action": "refuse", "scores": scores, "recovery": recovery}

```

Enhanced Audit Log

```

import hashlib
from datetime import datetime
class ARK6AuditLog:
    """
    Audit log with CPV vectors, ORIC scores, and recovery paths.
    """
    def __init__(self, config):
        self.config = config
        self.entries = []

    def record(self, output: str, gate_result: Dict, history: List[dict], context: str):
        """
        Log emission gate decision with dimensional details.
        """
        entry = {
            "timestamp": datetime.utcnow().isoformat(),

```



```

"output_hash": hashlib.sha256(output.encode()).hexdigest(),
"action": gate_result["action"],
"scores": gate_result["scores"],
"recovery": gate_result.get("recovery"),
"context_hash": hashlib.sha256(context.encode()).hexdigest(),
"history_size": len(history)
}

```

```

if self.config.audit.log_cpv_vectors and gate_result.get("recovery"):
    # Log CPV details if salvage occurred
    if "salvage" in gate_result["recovery"]:
        entry["cpv_fragments"] = gate_result["recovery"]["salvage"].get("fragments", [])

```

```

if self.config.audit.log_oric_scores:
    # Log ORIC dimensional breakdown
    b_details = gate_result["scores"].get("details", {}).get("verify", {})
    if "oric_scores" in b_details:
        entry["oric"] = b_details["oric_scores"]

```

```

self.entries.append(entry)

```

```

def get_entries(self):
    return self.entries

```

Complete A.R.K. 6.5 Controller

```

class ARK6:

```

```

    """

```

A.R.K. v6.5 - Dimensional Verifier Implementation

```

    """

```

```

def __init__(self, config):
    self.config = config
    self.gate = DimensionalEmissionGate(config)
    self.audit = ARK6AuditLog(config)

```

```
def process(self, output: str, history: List[dict], context: str) -> Dict:
```

```
    """
```

```
    Process output through dimensional verifiers and emission gate.
```

```
    Returns decision dict with action and detailed scores.
```

```
    """
```

```
    result = self.gate.gate(output, history, context)
```

```
    # Log to audit
```

```
    self.audit.record(output, result, history, context)
```

```
    return result
```

```
def get_audit_log(self):
```

```
    """Return audit entries for analysis."""
```

```
    return self.audit.get_entries()
```

```
Calibration Harness (calibrate.py)
```

```
import numpy as np
```

```
from sklearn.isotonic import IsotonicRegression
```

```
from sklearn.metrics import brier_score_loss
```

```
class CalibrationHarness:
```

```
    """
```

```
    Calibration tool for fitting thresholds and weights from validation traces.
```

```
    Uses isotonic regression for probability calibration and Brier score for evaluation.
```

```
    """
```

```
    def __init__(self, config, validation_data: List[Dict]):
```

```
        self.config = config
```

```
        self.validation_data = validation_data # List of {"output": str, "history": [], "context": str,  
"human_label": float (0-1 coherence)}
```

```
    def calibrate_cpv(self):
```

```
        """
```

```
        Calibrate CPV normalization using validation set.
```

Fits isotonic regression to map raw CPV to calibrated probabilities.

```
"""
raw_cpvs = []
labels = []

scorer = CPVCoherenceScorer(self.xconfig)

for data in self.validation_data:
    raw_cpv = scorer.compute_cpv(data["output"], data["history"]) # Raw before normalization
    raw_cpvs.append(raw_cpv)
    labels.append(data["human_label"])

ir = IsotonicRegression(out_of_bounds="clip")
ir.fit(raw_cpvs, labels)

# Update config or save model for runtime use
print("Calibrated CPV model fitted. Brier score:", brier_score_loss(labels, ir.predict(raw_cpvs)))
return ir
```

def optimize_thresholds(self, initial_thresholds: Dict) -> Dict:

```
"""
Optimize thresholds (emission, regen, refuse) via grid search on validation F1 or similar.
"""

# Placeholder grid search; in practice, use optuna or similar
best_thresholds = initial_thresholds
best_score = 0.0

gate = DimensionalEmissionGate(self.xconfig)

for emission in np.linspace(0.7, 0.8, 3):
    for regen in np.linspace(0.5, 0.6, 3):
        for refuse in np.linspace(0.35, 0.45, 3):
            self.xconfig.thresholds.emission = emission
```

```
self.config.thresholds.regen = regen
self.config.thresholds.refuse = refuse
```

```
predictions = []
for data in self.validation_data:
    result = gate.gate(data["output"], data["history"], data["context"])
    pred = 1 if result["action"] == "emit" else 0 # Binary for simplicity
    predictions.append(pred)
```

```
# Compute score (e.g., correlation with labels)
score = np.corrcoef(predictions, [d["human_label"] > 0.5 for d in self.validation_data])
```

```
[0,1]
if score > best_score:
    best_score = score
    best_thresholds = {"emission": emission, "regen": regen, "refuse": refuse}
```

```
print("Optimized thresholds:", best_thresholds)
return best_thresholds
```

Usage example

```
config = ... (load)
val_data = [{"output": "...", "history": [...], "context": "...", "human_label": 0.8}, ...]
harness = CalibrationHarness(config, val_data)
cpv_model = harness.calibrate_cpv()
thresholds = harness.optimize_thresholds(config.thresholds)
```

Usage Example

Load config

```
from types import SimpleNamespace
import yaml
with open("ark6_config.yaml") as f:
```

```
config_dict = yaml.safe_load(f)
```

Convert to nested SimpleNamespace for attribute access

```
def dict_to_namespace(d):
    if isinstance(d, dict):
```

```
return SimpleNamespace(**{k: dict_to_namespace(v) for k, v in d.items()})
```

```
return d
```

```
config = dict_to_namespace(config_dict)
```

```
Initialize A.R.K. 6.5
```

```
ark = ARK6(config)
```

```
Process an output
```

```
history = [
```

```
{ "text": "The meeting is tomorrow at 2pm." },
```

```
{ "text": "Should I bring the quarterly report?" }
```

```
]
```

```
context = "Planning for quarterly review meeting"
```

```
output = "Yes, bring the Q4 report. The meeting is at 3pm tomorrow."
```

```
result = ark.process(output, history, context)
```

```
print(f" Action: {result['action']}")
```

```
print(f" Coherence: {result['scores']['coherence']:.3f}")
```

```
print(f" Truth Density: {result['scores']['details']['trace']:.3f}")
```

```
print(f" Binding Consistency: {result['scores']['details']['verify']:.3f}")
```

```
print(f" Stability: {result['scores']['details']['engage']:.3f}")
```

```
if result.get('recovery'):
```

```
print(f" Recovery performed: {result['recovery']}")
```

```
Next Steps for Validation
```

- 1 Collect validation dataset: 1000 outputs with human quality labels
- 2 Run comparative benchmark: A.R.K. 6.5 vs standard NLI+retrieval
- 3 Measure:
 - Correlation with human judgments
 - False positive/negative rates
 - Recovery utility (salvaged outputs that humans approve)
- 4 Iterate on verifier implementations based on results
- 5 Calibrate thresholds/weights: Use calibrate.py on validation traces for empirical fitting

```
This is A.R.K. 6.5
```